TEF-RAG: Temporal Evidence Flow Retrieval-Augmented Generation for Energy Storage Station Operation and Maintenance

Juxian Yin

***Abstract*—An embodied operation and maintenance (O&M) system for an energy-storage station must consult records from different stages of a maintenance process. Alarms and inspections document the observed condition; work orders and procedures specify actions; post-action checks record their outcomes. Retrieval-augmented generation (RAG) supplies these records as context for work-order and action-plan generation. For a query issued at a given time, the retrieved records must already have been available and must jointly supply the evidence needed for diagnosis, action, and verification. Otherwise, the system may use retrospective information or provide an incomplete decision context. To address these requirements, we propose TEF-RAG (Temporal Evidence Flow Retrieval-Augmented Generation), which formulates retrieval in dynamic O&M environments as query-conditioned temporal evidence flow selection. TEF-RAG first applies deterministic constraints on asset identity, event time, availability time, and procedure applicability. A learned pair proposal module then selects evidence pairs worth semantic relation reasoning, and a directed graph is constructed with relations such as support, update, qualification, supersession, and verification. Finally, candidate evidence sets are constructed from the resulting graph, and a nonlinear RankNet learns their set-level relative utility to select an evidence flow of at most five records with complementary roles and coherent relations. To evaluate these temporal and cross-record requirements under controlled conditions, we construct a public-source-grounded synthetic temporal O&M benchmark with dual timestamps, supersession, cross-source complementary evidence, and persistent uncertainty, and compare different retrieval contexts under a shared downstream generator. On this benchmark, TEF-RAG records the highest Recall@5, Complete@5, and FlowComplete@5 among the six compared methods. Under the shared generator, TEF-RAG records the highest Field Macro F1 score, Action F1 score, Citation F1 score, and Evidence Support Recall. Thus, the comparison distinguishes individual-record ranking quality from the recovery of temporally valid and process-complete evidence sets.**

***Index Terms*—retrieval-augmented generation, temporal retrieval, evidence flow, learning to rank, operation and maintenance, energy storage**

I. Introduction

In embodied operation and maintenance (O&M) for energy-storage stations, decision making depends on dynamic and process-oriented knowledge. Handling a single anomaly often requires jointly consulting equipment states, alarms and inspection records, historical work orders, operating procedures, and post-action verification results. Unlike static knowledge question answering, the system must answer two stricter questions in addition to identifying which evidence is relevant to the query: whether that evidence was already available at the query time, and whether the evidence can jointly form a complete process that supports the current decision. These requirements make retrieval-based grounding useful for O&M generation because the generator must consult external records that change over time rather than rely only on knowledge stored in model parameters. Retrieval-augmented generation (RAG) combines language models with external retrieval for knowledge-intensive generation [4]. On the retrieval side, Best Matching 25 (BM25) remains a strong lexical baseline, while neural retrievers such as DPR and ColBERT improve semantic matching through dense representations or late interaction [5], [6]; BEIR further shows that retrieval paradigms differ substantially in cross-domain generalization [7]. On the generation side, FiD demonstrates that a generator can effectively aggregate multiple retrieved passages [8], while Self-RAG further studies joint control over retrieval, generation, and self-reflection [15]. However, conventional RAG typically treats candidate documents as independent relevance units. This assumption is insufficient for continuously updated O&M records. First, O&M evidence has dual temporal attributes: the event time and the time at which the record becomes available to the system need not coincide, and late-arriving records must not be used retrospectively by an earlier query. Second, maintenance decisions usually depend on complementary evidence: a state observation motivates an inspection, the inspection supports a diagnosis, the diagnosis leads to an action, and a post-action check verifies the preceding decision. High scores for individual documents therefore do not guarantee that the Top-k set covers the complete process.

Existing work addresses related aspects of this problem. Time-aware RAG models temporal constraints or information changes across time intervals [17], [18]; multi-hop retrieval conditions later retrieval steps on previously retrieved information [10], [11]; and RAPTOR, GraphRAG, and HippoRAG organize information across documents using hierarchical or graph structures [12]–[14]. These lines of work address temporal filtering, iterative retrieval, and cross-document organization from different angles. None of the cited methods explicitly combines query-time evidence visibility with direct set-level optimization for O&M process completeness. This gap motivates a retrieval objective that treats both requirements explicitly. We propose TEF-RAG (Temporal Evidence Flow Retrieval-Augmented Generation). Its core idea is to construct query-conditioned directed semantic relations among evidence that is visible at the query time, and to reformulate retrieval

from independent node ranking into evidence-set ranking. The system first removes unusable records with deterministic asset, temporal, and procedure constraints; it then learns which candidate evidence pairs are worth relation reasoning and constructs a temporal evidence graph; finally, it applies nonlinear pairwise ranking over size-bounded candidate sets so that the model learns which candidate set better satisfies the annotated completeness objective for the current query.

TEF-RAG differs from general graph retrieval in that its graph is not predefined by static entity relations, but is formed dynamically from the currently visible O&M evidence under the query. It differs from time-aware ranking in that temporal constraints define only the legality boundary, while the final objective remains the recovery of a role-complementary and relation-consistent evidence set. It also differs from stepwise multi-hop retrieval because the target structure may include branching, updates, supersession, and preserved uncertainty rather than being restricted to a single linear path. In the embodied O&M setting considered in this work, retrieval is not the endpoint. The selected evidence flow is passed to a shared generator that produces a standardized work order and a dependency-aware action plan, providing an evidence-grounded interface for downstream planners, constraint checkers, and robotic execution. This paper focuses on evidence retrieval and the structured planning context in this pipeline, and does not claim to solve closed-loop robotic control or field safety certification.

Our contributions are threefold:

1) We formulate temporal evidence flow retrieval for dynamic O&M knowledge, unifying temporal visibility and evidence-set completeness within a single retrieval framework.
2) We propose a method that combines learned evidence-pair proposal, query-conditioned relation graphs, and nonlinear set ranking, shifting the retrieval objective from independent document relevance to set-level process utility.
3) We construct a controlled temporal O&M benchmark and evaluate retrieval together with downstream structured generation under a shared generator, using both retrieval and evidence-grounded generation metrics.

II. Related Work

Retrieval-Augmented Generation and Neural Retrieval. Because TEF-RAG operates between candidate retrieval and downstream generation, both neural retrieval and multi-evidence generation provide direct points of comparison. Classical RAG combines an external retriever with a generator [4]; DPR and ColBERT represent dense dual-encoder retrieval and late-interaction retrieval, respectively [5], [6]. FiD demonstrates multi-evidence aggregation from the generation side [8], while BEIR systematically compares the generalization of sparse, dense, and reranking paradigms [7]. DPR and ColBERT rank passages individually, while FiD aggregates multiple retrieved passages at generation time. Our retrieval objective instead assigns utility to an evidence set and its temporal and semantic relations.

Time-Aware and Multi-Hop Retrieval. Time-aware retrieval and multi-hop retrieval address two complementary aspects of multi-evidence retrieval. Time-aware methods determine which evidence is relevant under a given temporal condition, while multi-hop methods retrieve multiple pieces of evidence whose usefulness depends on one another. MRAG and TA-RAG explicitly incorporate temporal conditions into retrieval [17], [18], whereas MDR and IRCoT retrieve complementary evidence through sequential retrieval or interleaved retrieval and reasoning [10], [11]. TEF-RAG shares both concerns: it enforces temporal validity and seeks multiple complementary records for the same task. It further distinguishes event time from availability time and models explicit relations among the retrieved evidence rather than treating the evidence as a simple linear chain.

Structured and Graph-Based RAG. A complementary line of work changes how retrieved knowledge is organized across documents rather than focusing primarily on temporal filtering or iterative retrieval. RAPTOR, GraphRAG, and HippoRAG organize cross-document information using hierarchical or graph-based structures [12]–[14]. Unlike methods that typically construct graphs from document hierarchies or entity relations, TEF-RAG defines edges jointly from the current query and the currently visible O&M records with rich process semantics.

Embodied O&M and Battery Intelligence. At the application level, RAG has also begun to support embodied field robotics and professional equipment operation. SayComply studies retrieval of operational-compliance knowledge for field robots, while RAGLRO applies retrieval-augmented language models to robotic operation tasks [19], [20]. In the battery domain, LiBrain explores combinations of time-series information, retrieved knowledge, and language models for battery diagnosis and management [21]. These studies connect retrieved operational knowledge or time-series information to robotic or battery-management tasks. Our work focuses on the preceding evidence-selection step so that downstream work orders and action plans use evidence that is visible at the query time and covers the required process roles.

The preceding methods determine what evidence is retrieved or how it is organized; TEF-RAG additionally requires a final mechanism for comparing alternative evidence sets. **Learning to Rank.** The final selection stage of TEF-RAG is a query-level learning-to-rank problem. RankNet learns relative ranking functions from pairwise preferences [16], while work such as ANCE demonstrates the importance of hard negatives for learning effective retrieval boundaries [9]. We apply this idea to evidence sets rather than individual documents: the training objective directly compares complete sets with high-scoring but incomplete hard negatives under the same query.

III. Problem Definition

A. Queries and Temporally Visible Evidence

We represent an O&M query as

$$Q_q = (x_q, u_q, m_q, t_q), \tag{1}$$

where x_q is the query text, u_q is the target asset, m_q denotes the asset-model/context metadata used for scope checks, and t_q is the knowledge cutoff time. The global evidence repository is denoted by $\mathcal{K} = \{e_1, \ldots, e_N\}$. Each evidence record contains at least its content, asset identifier, event type, source type, event time $t_e(e)$, and availability time $t_a(e)$.

We explicitly distinguish between when an event happened and when the system knew about it. Evidence e is visible to query q if and only if

$$\mathrm{Visible}(e, q) = \mathbb{I}[t_e(e) \le t_q \wedge t_a(e) \le t_q]\,. \tag{2}$$

Thus, a record describing an earlier event but entered into the system only after the query time cannot be used by that query.

For procedure evidence, we additionally require a matching asset model, validity at the query time, and that the procedure has not been withdrawn. We write these deterministic conditions collectively as

$$\mathrm{Legal}(e, Q_q) = \mathrm{Visible}(e, q) \wedge \mathrm{ScopeOK}(e, Q_q), \tag{3}$$

where $\mathrm{ScopeOK}$ summarizes asset consistency and, when applicable, procedure model scope, validity interval, and withdrawal constraints. The legal evidence set at query time is therefore

$$\mathcal{C}_q = \{e \in \mathcal{K} : \mathrm{Legal}(e, Q_q) = 1\}. \tag{4}$$

All subsequent retrieval, relation reasoning, and set ranking operate only within $\mathcal{C}_q$.

To distinguish the levels used below, an *evidence record* refers to one retrievable item, an *evidence set* refers to a selected collection of records, and a *temporal evidence flow* refers to such a set together with the directed relations among its records.

B. Retrieval Objective

A base retrieval stage produces a candidate node pool $V_q \subseteq \mathcal{C}_q$ from the legal evidence. For a pool V_q, the selector compares sets of $k_q = \min(5, |V_q|)$ records, with $k = 5$ as the retrieval upper bound. When fewer than five legal records exist, it returns those records without padding or backfilling. Within this budget, a useful set must cover the roles needed by the query and retain consistent temporal and semantic relations among its records.

The ranker scores each candidate set jointly:

$$\hat{\mathcal{S}}_q = \arg\max_{\mathcal{S} \in \mathcal{B}_q} F_\theta(Q_q, \mathcal{S}, \mathcal{G}_q), \quad |\mathcal{S}| = k_q, \tag{5}$$

Here F_θ ranks the sets of size k_q in $\mathcal{B}_q$, using the query-conditioned temporal evidence flow graph $\mathcal{G}_q$.

TABLE I
Main notation.

Symbol	Meaning
$Q_q = (x_q, u_q, m_q, t_q)$	Query text, asset, asset-model/context metadata, and knowledge cutoff time
$e, t_e(e), t_a(e)$	Evidence record, event time, and first availability time
$\mathcal{K}$	Global evidence repository
$\mathcal{C}_q$	Legal and visible evidence set at query time
V_q	Candidate node pool returned by the base retrieval stage
$\mathcal{G}_q$	Query-conditioned temporal evidence flow graph
$\mathcal{B}_q$	Candidate space of evidence sets of size $k_q \le 5$
$\hat{\mathcal{S}}_q$	Final selected evidence set (at most 5 records)
W_q, A_q	Standardized work order and action plan

C. Downstream Task

The retrieval result is used as controlled context for a structured generator, which produces a standardized work order W_q and an action plan A_q:

$$(\hat{W}_q, \hat{A}_q) = g_\psi(x_q, \hat{\mathcal{S}}_q). \tag{6}$$

The evidence records in $\hat{\mathcal{S}}_q$ are passed to the generator as the retrieved context; retrieval scores and temporal evidence flow graph features are not exposed to the generator. The primary algorithmic contribution of this paper is the retrieval stage that constructs and ranks temporal evidence flows. Structured generation experiments evaluate whether differences in retrieved evidence are reflected in work-order fields, actions, dependencies, and evidence citations.

IV. TEF-RAG Method

Figure 1 presents the overall TEF-RAG pipeline. The method consists of five connected stages: legal-candidate retrieval, learned evidence-pair proposal, query-conditioned relation graph construction, candidate evidence-set construction, and nonlinear set ranking. The first two stages control the candidate and relation-judging budgets, the relation graph represents process structure among evidence records, the fourth stage constructs alternative evidence sets, and the final ranker compares their overall utility.

A. Candidate Evidence Preparation and Temporal Legality

TEF-RAG first applies a unified legality check to heterogeneous raw records. Records are directly excluded if they refer to the wrong asset, describe events after the query time, become available only after the query time, or correspond to procedures that are inapplicable to the current model or time period. This stage does not learn relevance; instead, it establishes a knowledge boundary that subsequent models cannot override.

Within the legal set $\mathcal{C}_q$, BM25 produces a broad candidate pool. The candidate records are assigned an initial node score

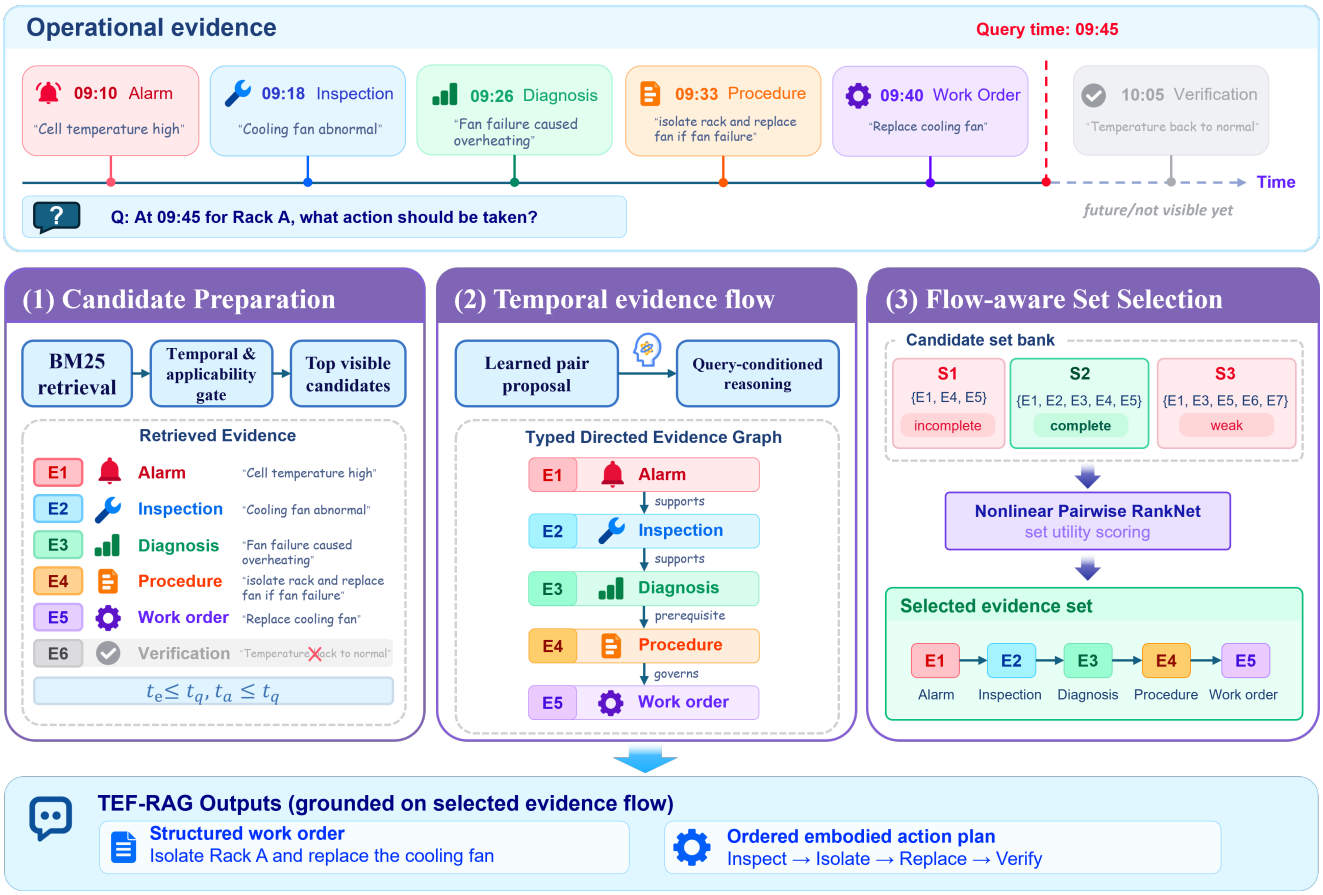


Fig. 1. Overview of TEF-RAG. The system first selects candidate evidence pairs from records that are legal at the query time, then constructs a query-conditioned temporal evidence flow. Candidate evidence sets are constructed from the resulting retrieval and graph signals, and a nonlinear pairwise ranker selects at most five records as the final evidence set for a shared downstream generator.

based on lexical relevance, temporal recency, and compatibility with the evidence roles implied by the query. The 30 highest-scoring legal records form the candidate node pool V_q defined in Section III. This pool is designed to preserve potentially useful evidence for subsequent relation reasoning and evidence-set selection rather than to determine the final output directly.

B. Learned Pair Proposal

Applying semantic relation reasoning to every candidate pair in the search pool would incur quadratic cost in the number of candidates. TEF-RAG therefore first learns a lightweight pair proposer that estimates, for each temporally valid directed pair (e_i, e_j), the probability that the pair is worth passing to the relation-reasoning stage:

$$p_{ij} = \sigma\left(\mathbf{w}^\top \phi_p(q, e_i, e_j) + b\right). \tag{7}$$

Here, $\phi_p(q, e_i, e_j)$ denotes the feature representation of query q and evidence pair (e_i, e_j). It includes textual relevance between the query and both records, similarity between the two records, event types, temporal gaps, and supersession-related information. If an evidence pair realizes a required relation in an annotated reference evidence flow from the training data, it is treated as a positive sample; all other temporally valid pairs are treated as negative samples. Logistic regression is then used to learn the priority of each candidate pair, and only high-scoring pairs are forwarded to the subsequent relation-reasoning module.

In the configuration used for evaluation, a Top-30 search pool contains at most $\binom{30}{2} = 435$ temporally valid pairs after imposing a single forward temporal direction, whereas the pair proposer retains only the 32 highest-scoring pairs for subsequent relation judging. Thus, when the search pool contains 30 records, the number of relation judgments is reduced by up to approximately 92.6%. The primary role of pair proposal is computational budget control: it sends only the highest-scoring candidate pairs to semantic relation reasoning while preserving the deterministic temporal direction constraint.

C. Query-Conditioned Temporal Evidence Flow

For each retained candidate pair, TEF-RAG uses a query-conditioned large language model (LLM) relation judge that reads the query together with the text and metadata of both evidence records and outputs

$$h_\eta(q, e_i, e_j) \rightarrow (r_{ij}, c_{ij}), \tag{8}$$

where $r_{ij} \in \mathcal{R}$ and $c_{ij} \in [0,1]$ is the relation confidence. Only pairs predicted to have a valid relation and exceeding the confidence threshold are added to the graph. Pairs assigned no valid relation or a confidence below the threshold are discarded. The fixed prompt used by the relation judge is provided in Appendix A-C.

Candidate evidence records are treated as graph nodes, and identified semantic relations between two records are represented as directed edges. For query q, the temporal evidence flow graph is

$$\mathcal{G}_q = (V_q, E_q), \tag{9}$$

where V_q is the candidate node pool and E_q is the set of accepted evidence relations. We write an edge as (e_i, r, e_j): e_i is the source record, e_j the target, and $r \in \mathcal{R}$ their relation.

Edges take one of ten relation types: supports, governs, qualifies, contrasts, prerequisite, preserves uncertainty, resolves, supersession, updates, and verifies. Event time or an explicit prerequisite/update relation determines the direction. Branches and merges are allowed.

The query determines how a pair of records is interpreted in $\mathcal{G}_q$; its task role can change with the information requested. The relation judge also cannot override the legality boundary defined in Section III; thus, late-arriving evidence, future events, and inapplicable procedures cannot re-enter the candidate set through semantic reasoning.

D. Candidate Evidence-Set Construction

After the relation graph is built, TEF-RAG compares different evidence combinations rather than continuing to rank individual records independently. Enumerating all five-record combinations from the full search pool is expensive, so we restrict ranking to candidate sets generated by the procedures below.

Specifically, we enumerate five-record combinations from the highest-scoring evidence nodes to cover different combinations of highly relevant records. We also retain sets produced by the preceding retrieval steps and perform single-record replacements around them, allowing lower-ranked evidence that fills a required task role to enter the candidate-set collection.

For each set $\mathcal{S}$, we extract a set-level representation $\phi_s(q, \mathcal{S}, \mathcal{G}_q)$. The frozen feature schema used by the final model contains 62 dimensions. The 62 features fall into the following groups:

- node-score statistics for relevance, recency, role compatibility, and the aggregate score;
- lexical similarity between records, spans of event and availability times, and event-gap statistics;
- counts and diversity of event types and source types;
- coverage of required task roles, counts of unmet demands, and uncertainty indicators for the query and set;
- graph measures: active edges, edge confidence and density, node degrees, connected components, and directed-path statistics;
- the node, edge, role-coverage, redundancy, connectivity, and uncertainty-consistency terms from the rule-based score used to select first-round hard negatives.

Role coverage and graph features let the ranker distinguish sets with similar node scores but different missing roles or internal relations.

E. Nonlinear Pairwise Set Ranking

The final ranker is a nonlinear RankNet [16]. Let

$$s_\theta(q, \mathcal{S}) = \mathrm{MLP}_\theta(\phi_s(q, \mathcal{S}, \mathcal{G}_q)). \tag{10}$$

The final stage uses pairwise preference learning over evidence sets. For the same query, the candidate-set collection typically contains multiple complete and incomplete alternatives. Each training pair $(\mathcal{S}_i^+, \mathcal{S}_j^-)$ compares a set $\mathcal{S}_i^+$ that forms a complete reference evidence flow with a set $\mathcal{S}_j^-$ that lacks required records or relations.

For this pair, the RankNet loss favors $\mathcal{S}_i^+$:

$$\ell(q, \mathcal{S}_i^+, \mathcal{S}_j^-) = -\log \sigma\left(s_\theta(q, \mathcal{S}_i^+) - s_\theta(q, \mathcal{S}_j^-)\right). \tag{11}$$

Training minimizes the mean pairwise loss:

$$\mathcal{L} = \frac{1}{|\mathcal{P}|} \sum_{(q, \mathcal{S}_i^+, \mathcal{S}_j^-) \in \mathcal{P}} \ell(q, \mathcal{S}_i^+, \mathcal{S}_j^-), \tag{12}$$

Here $\mathcal{P}$ contains the training pairs defined above.

The first of two training rounds draws hard negatives using a rule-based set score (Appendix A-B). The score accounts for evidence quality, relations within the set, role coverage, redundancy, and isolated records. An incomplete set can score well when its records are individually relevant, even though a required evidence category or relation is missing. We prioritize these sets over random negatives and retain a small number of other set compositions for diversity.

For the second round, we use the trained ranker to rescore all negative sets and select those it ranks highest. These are the incomplete flows that the first-round model mistakes for useful sets. Training on them penalizes those ranking errors directly.

F. Interface to the Embodied O&M Generator

The output of TEF-RAG is not a low-level robotic control command, but an evidence context for downstream embodied execution. A shared downstream generator takes the query and retrieved evidence to produce a standardized work order and a dependency-aware action plan; these structured outputs can then be passed to a planner, constraint checker, or robotic executor.

To isolate the effect of retrieval quality, all compared methods share the same downstream generator. The generator receives only the query and the evidence records selected by the current retrieval method; it does not receive the method name, retrieval scores, or TEF graph structure. The fixed generation prompt is provided in Appendix A-D. Our evaluation stops at the structured work order and action plan, and does not treat closed-loop robotic execution or field safety certification as solved problems.

TABLE II
DATA HIERARCHY OF THE BENCHMARK.

Level	Meaning
Chain	A complete O&M event process
Evidence	An atomic retrievable record with event time and availability time
Intent	A distinct O&M information need defined on the same chain
Query	A differently worded query expressing the same intent
Reference flow	Key evidence requirements and their relation constraints

V. EXPERIMENTAL DESIGN

A. Temporal O&M Benchmark

We construct a public-source-grounded synthetic temporal O&M benchmark for controlled evaluation of temporal evidence retrieval in dynamic O&M scenarios. The dataset is constrained by public-source background information and constructed with artificial intelligence (AI) assistance; it is not derived from real power-station fault logs and is not a real O&M dataset manually annotated by domain experts. We use public product documents and research datasets to constrain cell type, typical physical ranges, and sampling background. The O&M events, diagnostic changes, procedure contexts, and queries are then synthesized under these constraints. We include late-arriving records, supersession, cross-source complementary evidence, multi-episode ambiguity, and persistent uncertainty as controlled conditions for comparing retrieval methods.

Figure 2 traces the construction from public-source constraints to synthetic O&M event chains. Each chain supplies the records to be retrieved and the events about which queries are asked. We annotate each query with a reference evidence flow and a structured reference answer, used for retrieval and generation evaluation, respectively.

Table II lists the levels of the dataset. A chain contains the records of a complete O&M process. Several task intents are defined on that chain, and two differently worded queries express each intent.

The reference annotations specify the information a query requires and the relations needed among its supporting records. Different records may satisfy the same requirement. Evaluation checks these requirements and relations across the retrieved set, so equivalent evidence can receive credit.

Across 100 target assets, the benchmark comprises 400 O&M event chains with 3,888 retrievable records. These chains yield 1,200 task intents and 2,400 queries. Development, validation, and test sets are split at the complete event-chain level. All queries derived from the same event chain appear in only one data split, preventing the same event content from appearing in both training and test data through different phrasings.

To evaluate retrieval under both ordinary and deliberately difficult temporal conditions, we divide the 400 event chains evenly into two difficulty strata. The standard subset contains relatively common combinations of temporal and O&M conditions. The temporal-hard subset increases the frequency of cases involving late-arriving evidence, supersession, multi-episode ambiguity, cross-source complementary evidence, and persistent uncertainty.

Public sources are used only to constrain cell type, typical operating ranges, and temporal sampling background, including energy-storage cell documents from HiTHIUM and REPT BATTERO and the RWTH M5BAT dataset [22]–[25]. These sources are not used to define synthetic fault frequencies, station-level thresholds, or specific maintenance causal narratives in our benchmark.

B. Structured Generation Evaluation

For generation evaluation, we prepare structured reference answers from the same O&M event chains. Each answer is checked repeatedly against its chain and the required output format before it is finalized.

The two queries for an intent ask for the same diagnosis, action, and verification information, so both are evaluated against one reference answer. There are 720 reference answers for the 1,440 development queries, 240 for the 480 validation queries, and 240 for the 480 test queries. Each output has two parts: a standardized work order and an action plan. The work order describes the diagnosis, recommended actions, applicable procedure, verification/uncertainty status, and evidence citations; the action plan consists of atomic actions and their dependencies.

C. Baselines

To reduce confounding from differences outside the retrieval strategy, all methods first operate on the same evidence set that is eligible at the query time. We retain only records for the target asset satisfying $t_e(e) \leq t_q$ and $t_a(e) \leq t_q$, and apply the same version-validity, withdrawal-status, and model-applicability constraints to procedure evidence. We compare six methods under this shared setup, with a final retrieval budget of at most five evidence records for every method.

BM25. We use standard BM25 as the lexical retrieval baseline [1]. For each query, a separate index is built over the eligible evidence set, with fixed parameters $k_1 = 1.5$ and $b = 0.75$, and the five highest-scoring evidence records are returned directly.

BGE Reranker. We first retrieve the Top-30 candidates using the same BM25 setup and then use `BAAI/bge-reranker-v2-m3` to rescore each $(\text{query}, \text{evidence})$ text pair [2], [3]. The top five records under the reranking score are selected. This baseline tests whether a general-purpose semantic reranker can compensate for the limitations of purely lexical retrieval.

Fine-tuned BGE Reranker. To control for task-specific supervision, we additionally adapt the same BGE reranker using only development-split document-level relevance labels. Evidence records listed as acceptable for any required evidence group are treated as positives, while high-ranked

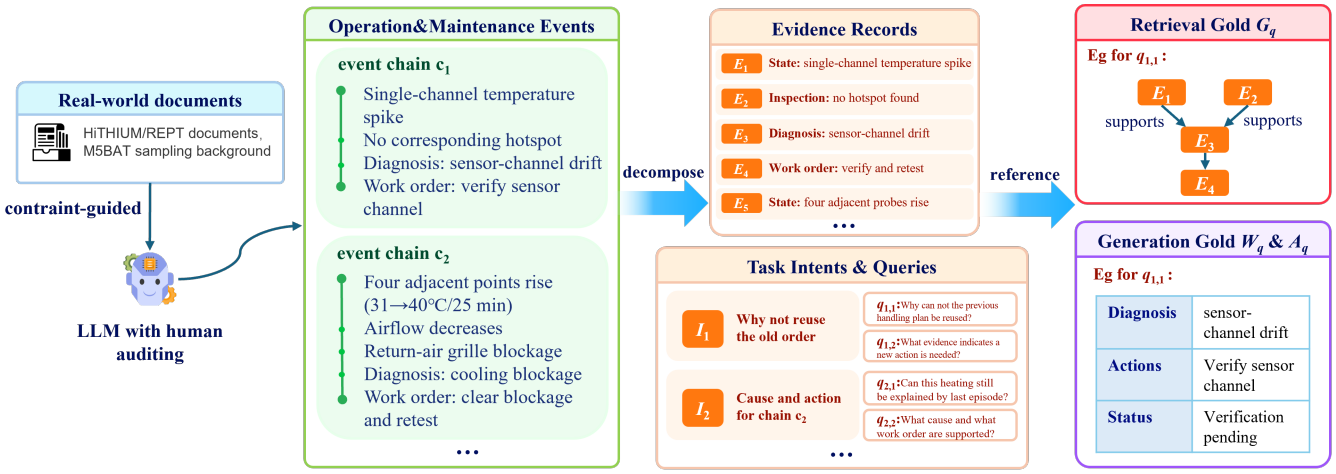


Fig. 2. Construction pipeline and example for the temporal O&M benchmark. Public sources constrain the equipment and data background; an O&M event chain c is then constructed and used to derive evidence records e_i, task intents I_j, and queries $q_{j,p}$. Retrieval reference evidence flows and structured generation reference answers are subsequently created for each query to support retrieval and downstream generation evaluation.

TABLE III
SCALE AND SPLIT OF THE TEMPORAL O&M BENCHMARK.

Statistic	Development	Validation	Test	Total
Authored chains	240	80	80	400
Primary intents	720	240	240	1200
Query rows	1440	480	480	2400
Evidence records	2335	779	774	3888

eligible BM25 nonpositives are used as negatives; no relation or evidence-flow labels are used. We train rank-8 low-rank adapters on the attention query/value projections together with the classification head for two epochs, and select the checkpoint by validation normalized discounted cumulative gain at 5 (nDCG@5), with Recall@5 as a tie-breaker. At inference, this baseline uses the same BM25 Top-30 candidate pool and Top-5 output budget as the pretrained BGE reranker.

Temporal-BM25. This baseline explicitly adds temporal matching to BM25 relevance. Its ranking score is

$$s_{\mathrm{TBM25}}(e,q) = \lambda\, \widetilde{s}_{\mathrm{BM25}}(e,q) + (1-\lambda)\, s_{\mathrm{time}}(e,q), \quad (13)$$

where $\widetilde{s}_{\mathrm{BM25}}$ is the normalized BM25 score and s_{time} measures the match between the evidence event time and the target time implied by expressions such as years, time intervals, before/after, and current/latest in the query. For queries without explicit temporal intent, the temporal term is set to a constant, preserving the original BM25 ranking. Before final test evaluation, we compare $\lambda \in \{0.25, 0.50, 0.75\}$ on validation data, select $\lambda = 0.25$, and fix this value for subsequent evaluation.

TA-RAG. We adopt the Time-Aware RAG method of Lau et al. [18]. The method first decomposes a query into its topic and target time range, then restricts candidate documents using temporal-interval filtering and performs semantic retrieval and BGE reranking within the temporal candidate set. To adapt TA-RAG to our discrete O&M events, each event time $t_e(e)$ is represented as a one-second half-open interval, and relative temporal expressions in a query are interpreted with t_q as the reference time. In addition, the TA-RAG index is built only from evidence that satisfies the same temporal and procedure-applicability constraints as the other methods, preventing future records from affecting retrieval. Apart from substituting the language model used for temporal decomposition, the temporal decomposition, interval filtering, Nomic embedding, FAISS retrieval, and BGE reranking pipeline follows the public implementation.

TEF-RAG. Our method applies learned evidence-pair screening, query-conditioned relation-graph construction, and nonlinear set-level ranking to the same eligible evidence set, and returns the candidate set with the highest overall evidence-flow score.

D. Evaluation Metrics

Retrieval. Recall@5 and nDCG@5 measure conventional retrieval relevance. Complete@5 measures whether the retrieved set contains all categories of key evidence required to complete the task. FlowComplete@5 further checks whether those records can form correct and mutually consistent evidence relations, thereby constituting a complete temporal evidence flow.

Generation. We report four primary metrics in the main text: Field Macro F1 (strict), Action F1, Citation F1, and Evidence Support Recall. For the F1 measures, F1 denotes the harmonic mean of precision and recall. Field Macro F1

TABLE IV
MAIN RETRIEVAL RESULTS ON THE TEST SPLIT (TOP-5).

Method	Recall@5	nDCG@5	Complete@5	FlowComplete@5
BM25	0.3791	0.3915	0.0563	0.0563
BGE Reranker	0.2916	0.3070	0.0417	0.0417
Fine-tuned BGE	0.6650	**0.6501**	0.2854	0.2833
Temporal-BM25	0.4591	0.4955	0.0792	0.0792
TA-RAG	0.2769	0.2907	0.0375	0.0375
TEF-RAG	**0.7228**	0.6231	**0.3188**	**0.3146**

measures strict matching of key work-order fields, Action F1 measures normalized maintenance actions, and Citation F1 evaluates claim–evidence citation pairs. Evidence Support Recall measures the fraction of reference work-order and action claims requiring evidence support for which the generated output provides at least one acceptable supporting-evidence citation.

E. Ablation Setup

To analyze the contribution of individual components, we define five ablation variants and evaluate their frozen retrieval outputs on the same 480-query test split used for the main retrieval comparison. The ablations remove learned pair proposal, relation-graph-derived set features, hard-negative mining, the nonlinear set ranker, or both learned pair proposal and the nonlinear set ranker. The full model and all variants use the same Top-30 search pool, the same maximum number of evidence pairs sent to relation judging, the same relation-judging prompt and confidence threshold, and the same candidate-set construction procedure. The graph-feature ablation removes only relation-graph-derived set features and retrains the same RankNet; the hard-negative ablation keeps only the first training round; and the nonlinear-ranker ablation removes the learned set ranker and uses the preceding greedy evidence-set selector instead. The joint ablation removes both learned components: the pair proposer is replaced by the nonlearned pair prefilter, and the nonlinear ranker is replaced by the greedy evidence-set selector. Relation judging, relation-graph construction, and candidate-set construction remain unchanged. For the test evaluation, predictions are frozen before the test evaluator is accessed.

VI. EXPERIMENTAL RESULTS

A. Main Retrieval Results

Table IV reports retrieval performance on the test split. TEF-RAG records the highest Recall@5, Complete@5, and FlowComplete@5, while the Fine-tuned BGE Reranker obtains the highest nDCG@5. Compared with Fine-tuned BGE, TEF-RAG increases Recall@5 from 0.6650 to 0.7228, Complete@5 from 0.2854 to 0.3188, and FlowComplete@5 from 0.2833 to 0.3146, whereas Fine-tuned BGE has higher nDCG@5 (0.6501 versus 0.6231). This pattern is consistent with the distinction between ranking individual evidence records and recovering a complete evidence set.

TABLE V
TEF-RAG ABLATION RESULTS ON THE TEST SPLIT.

Variant	Recall@5	nDCG@5	Complete@5	FlowComplete@5
Full TEF-RAG	0.7228	0.6231	0.3188	0.3146
w/o Pair Proposal	0.6829	0.5870	0.2667	0.2604
w/o Graph Features	0.6558	0.5723	0.2375	0.2354
w/o Hard Negatives	0.6540	0.5557	0.1979	0.1917
w/o Nonlinear Ranker	0.6242	0.6124	0.2042	0.1938
w/o Pair Proposal + Nonlinear Ranker	0.5835	0.6034	0.1625	0.1521

B. Test Ablation Study

Table V reports all five ablations on the same 480-query test split as the main retrieval comparison. Removing the learned pair proposer lowers Recall@5 from 0.7228 to 0.6829 and FlowComplete@5 from 0.3146 to 0.2604. Removing relation-graph-derived set features gives Recall@5 of 0.6558 and FlowComplete@5 of 0.2354. Ablating either learned pair selection or graph-derived set features leads to lower test-set performance.

Hard-negative mining and the nonlinear set ranker have the largest single-component effects on FlowComplete@5. Removing hard-negative mining lowers FlowComplete@5 from 0.3146 to 0.1917, while replacing the nonlinear ranker with the greedy evidence-set selector gives 0.1938. The latter variant also lowers nDCG@5 from 0.6231 to 0.6124, together with reductions in Recall@5 from 0.7228 to 0.6242 and Complete@5 from 0.3188 to 0.2042. On the test split, the nonlinear set ranker therefore improves both conventional ranking quality and evidence-set completeness.

Jointly removing the learned pair proposer and nonlinear ranker lowers Recall@5 to 0.5835, nDCG@5 to 0.6034, Complete@5 to 0.1625, and FlowComplete@5 to 0.1521. Even in this joint configuration, all four test metrics remain higher than BM25, the pretrained BGE Reranker, Temporal-BM25, and TA-RAG; Fine-tuned BGE and the full TEF-RAG remain higher on all four. Thus, temporal eligibility, relation reasoning, and candidate-set construction retain useful retrieval capability even after both benchmark-trained selection modules are removed.

C. Downstream Structured Generation Results

Table VI reports four complementary generation metrics used to compare structured outputs produced from each method's retrieved evidence. Under the same generator, prompt, output format, and decoding settings, TEF-RAG records the highest value on all four displayed metrics: Field Macro F1 of 0.5345, Action F1 of 0.2035, Citation F1 of 0.1585, and Evidence Support Recall of 0.2253. Fine-tuning BGE improves all four displayed metrics relative to the pretrained BGE reranker, but TEF-RAG remains higher on each. Under the shared generator, the TEF-RAG retrieval context therefore yields the strongest field-level, action-level, citation, and evidence-support results among the compared methods.

VII. LIMITATIONS AND THREATS TO VALIDITY

The benchmark in this work is a public-source-grounded, synthetic benchmark rather than a collection of real power-

TABLE VI
PRIMARY STRUCTURED-GENERATION METRICS ON THE GENERATION TEST SET.

Method	Field Macro F1	Action F1	Citation F1	Evidence Support Recall
BM25	0.4715	0.1372	0.0813	0.1258
BGE Reranker	0.4129	0.1198	0.0405	0.0614
Fine-tuned BGE	0.4818	0.1910	0.0954	0.1328
Temporal-BM25	0.5218	0.1637	0.1249	0.1871
TA-RAG	0.4059	0.1115	0.0379	0.0570
TEF-RAG	**0.5345**	**0.2035**	**0.1585**	**0.2253**

station work orders or domain-expert annotations. Its primary value therefore lies in controlled comparison of temporal evidence-retrieval mechanisms; the results should not be directly extrapolated to real-world power-station fault frequencies or field robotic safety. Although the generation evaluation covers fields, actions, dependencies, and citations, it does not simulate actual robotic execution, environmental feedback, or safety interlocks.

VIII. CONCLUSION

We presented TEF-RAG, which extends the retrieval objective for dynamic O&M RAG from independent document relevance to a temporal evidence flow that is visible at the query time, relation-consistent, and process-complete. On the test split, TEF-RAG records the highest Recall@5, Complete@5, and FlowComplete@5 among the six compared methods, while Fine-tuned BGE records the highest nDCG@5. In the structured generation test with a fixed downstream generator, TEF-RAG records the highest Field Macro F1, Action F1, Citation F1, and Evidence Support Recall. Even after removing both learned selection modules, the joint ablation still outperforms BM25, the pretrained BGE reranker, Temporal-BM25, and TA-RAG on all four test retrieval metrics. A next step is to include planning and execution constraints in the retrieval objective.

REFERENCES

[1] S. Robertson and H. Zaragoza, "The Probabilistic Relevance Framework: BM25 and Beyond," *Foundations and Trends in Information Retrieval*, vol. 3, no. 4, pp. 333–389, 2009, doi: 10.1561/1500000019.

[2] J. Chen, S. Xiao, P. Zhang, K. Luo, D. Lian, and Z. Liu, "M3-Embedding: Multi-Linguality, Multi-Functionality, Multi-Granularity Text Embeddings Through Self-Knowledge Distillation," in *Findings of the Association for Computational Linguistics: ACL 2024*, pp. 2318–2335, 2024, doi: 10.18653/v1/2024.findings-acl.137.

[3] FlagOpen, "FlagEmbedding: Retrieval and Retrieval-Augmented LLMs," GitHub repository, 2026. [Online]. Available: https://github.com/FlagOpen/FlagEmbedding. Accessed: Sep. 28, 2026.

[4] P. Lewis, E. Perez, A. Piktus, *et al.*, "Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks," in *Advances in Neural Information Processing Systems*, vol. 33, pp. 9459–9474, 2020.

[5] V. Karpukhin, B. Oguz, S. Min, *et al.*, "Dense Passage Retrieval for Open-Domain Question Answering," in *Proc. EMNLP*, pp. 6769–6781, 2020, doi: 10.18653/v1/2020.emnlp-main.550.

[6] O. Khattab and M. Zaharia, "ColBERT: Efficient and Effective Passage Search via Contextualized Late Interaction over BERT," in *Proc. ACM SIGIR*, pp. 39–48, 2020, doi: 10.1145/3397271.3401075.

[7] N. Thakur, N. Reimers, A. Rücklé, A. Srivastava, and I. Gurevych, "BEIR: A Heterogeneous Benchmark for Zero-shot Evaluation of Information Retrieval Models," in *Advances in Neural Information Processing Systems*, vol. 34, 2021.

[8] G. Izacard and E. Grave, "Leveraging Passage Retrieval with Generative Models for Open Domain Question Answering," in *Proc. EACL*, pp. 874–880, 2021, doi: 10.18653/v1/2021.eacl-main.74.

[9] L. Xiong, C. Xiong, Y. Li, *et al.*, "Approximate Nearest Neighbor Negative Contrastive Learning for Dense Text Retrieval," in *International Conference on Learning Representations*, 2021.

[10] W. Xiong, X. L. Li, S. Iyer, *et al.*, "Answering Complex Open-Domain Questions with Multi-Hop Dense Retrieval," in *International Conference on Learning Representations*, 2021.

[11] H. Trivedi, N. Balasubramanian, T. Khot, and A. Sabharwal, "Interleaving Retrieval with Chain-of-Thought Reasoning for Knowledge-Intensive Multi-Step Questions," in *Proc. ACL*, pp. 10014–10037, 2023, doi: 10.18653/v1/2023.acl-long.557.

[12] P. Sarthi, S. Abdullah, A. Tuli, S. Khanna, A. Goldie, and C. D. Manning, "RAPTOR: Recursive Abstractive Processing for Tree-Organized Retrieval," in *International Conference on Learning Representations*, 2024.

[13] D. Edge, H. Trinh, N. Cheng, *et al.*, "From Local to Global: A Graph RAG Approach to Query-Focused Summarization," *arXiv preprint arXiv:2404.16130*, 2024.

[14] B. Jiménez Gutiérrez, Y. Shu, Y. Gu, M. Yasunaga, and Y. Su, "HippoRAG: Neurobiologically Inspired Long-Term Memory for Large Language Models," in *Advances in Neural Information Processing Systems*, vol. 37, 2024, doi: 10.52202/079017-1902.

[15] A. Asai, Z. Wu, Y. Wang, A. Sil, and H. Hajishirzi, "Self-RAG: Learning to Retrieve, Generate, and Critique through Self-Reflection," in *International Conference on Learning Representations*, 2024.

[16] C. J. C. Burges, T. Shaked, E. Renshaw, A. Lazier, M. Deeds, N. Hamilton, and G. Hullender, "Learning to Rank Using Gradient Descent," in *Proc. ICML*, pp. 89–96, 2005, doi: 10.1145/1102351.1102363.

[17] S. Zhang, Y. Xue, Y. Zhang, X. Wu, A. T. Luu, and C. Zhao, "MRAG: A Modular Retrieval Framework for Time-Sensitive Question Answering," in *Findings of EMNLP*, pp. 3080–3118, 2025, doi: 10.18653/v1/2025.findings-emnlp.167.

[18] K. H. Lau, R. Zhang, W. Shi, X. Zhou, and X. Cheng, "Reading Between the Timelines: RAG for Answering Diachronic Questions," *arXiv preprint arXiv:2507.22917*, 2025.

[19] M. F. Ginting, D.-K. Kim, S.-K. Kim, *et al.*, "SayComply: Grounding Field Robotic Tasks in Operational Compliance through Retrieval-Based Language Models," in *Proc. IEEE ICRA*, pp. 13730–13736, 2025, doi: 10.1109/ICRA55743.2025.11128684.

[20] W. Wang, *et al.*, "RAGLRO: Retrieval-Augmented Generation With Large Language Models for Robotic Operations," *CAAI Transactions on Intelligence Technology*, vol. 11, no. 2, pp. 385–395, 2026, doi: 10.1049/cit2.70105.

[21] Z. Li, *et al.*, "LiBrain: LLM-Powered Li-ion Battery Diagnostics with Time-Series-Aware Retrieval-Augmented Framework for E-bikes," in *Proc. AAAI*, vol. 40, no. 47, pp. 40045–40053, 2026, doi: 10.1609/aaai.v40i47.41439.

[22] HiTHIUM, "HiTHIUM V1.1 ESS Cell 280 Ah EU English Datasheet," V1.1, 2023.

[23] HiTHIUM, "280 Ah LFP Cell Datasheet," V3.3, 2024.

[24] REPT BATTERO, "REPT BATTERO Battery Energy Storage Products Brochure (English Version)," 2025.

[25] S. Zurmühlen, L. Koltermann, and D. U. Sauer, "M5BAT Large-Scale Battery Storage System: Dataset for Battery Unit Pb1 2017–2025," RWTH Aachen University, 2026, doi: 10.18154/RWTH-2026-06637.

APPENDIX A
IMPLEMENTATION DETAILS AND MODEL PROMPTS

A. LLM Invocation Configuration

All LLM calls in this work, including temporal evidence flow relation judging, TA-RAG temporal decomposition, and downstream structured generation, directly use the official DeepSeek API with `deepseek-flash`. Apart from the frozen system/user prompts and input schemas for each task, sampling and decoding parameters are not manually overridden; all remaining API parameters follow the provider defaults.

B. Rule-Based Set Score for First-Round Hard Negatives

Because no learned set ranker is available before the first RankNet training round, we first use a fixed rule-based set score $H(\mathcal{S})$ to select harder negative samples from incomplete candidate sets. This score is used only for negative-sample selection and is not the final TEF-RAG prediction score. It is defined as

$$H(\mathcal{S}) = S_{\text{node}}(\mathcal{S}) + 0.48\, S_{\text{edge}}(\mathcal{S}) + 0.42\, S_{\text{role}}(\mathcal{S}) - 0.16\, S_{\text{red}}(\mathcal{S}) - 0.10\, S_{\text{disc}}(\mathcal{S}). \tag{14}$$

The node term is the sum of the evidence-node scores within the set:

$$S_{\text{node}}(\mathcal{S}) = \sum_{e \in \mathcal{S}} s_{\text{node}}(e), \tag{15}$$

where the score of an individual evidence record is

$$s_{\text{node}}(e) = 0.62\, s_{\text{rel}}(e) + 0.18\, s_{\text{rec}}(e) + 0.15\, s_{\text{role}}(e) + 0.05\, s_{\text{src}}(e). \tag{16}$$

Here, s_{rel} is the min–max-normalized BM25 score among the current legal candidates. The temporal-recency term is $s_{\text{rec}} = \exp(-\Delta t/T)$, with $T = 45$ days for ordinary queries and $T = 180$ days for queries involving comparisons across time. If the evidence type matches a task role required by the current query, $s_{\text{role}} = 1$; otherwise it is set to 0.35. For all legal candidates we set $s_{\text{src}} = 1$. It adds the same constant to every node score, leaving the ranking unchanged.

The relation term counts only valid relation edges internal to the candidate set. To avoid giving an unfair advantage to dense graphs simply because they contain more edges, only the highest-confidence incoming edge is retained for each target node:

$$S_{\text{edge}}(\mathcal{S}) = \sum_{e_j \in \mathcal{S}} \max_{(e_i, r, e_j) \in E(\mathcal{S})} c_{ij}, \tag{17}$$

where a node with no incoming edge contributes 0.

Role coverage is defined as

$$S_{\text{role}}(\mathcal{S}) = \frac{|R(\mathcal{S}) \cap D_q|}{|D_q|}, \tag{18}$$

where $R(\mathcal{S})$ is the set of evidence roles appearing in the candidate set and D_q is the set of roles required by the current query. The system considers state observation, diagnosis, and work order by default, and adds roles such as procedure, verification, inspection, uncertainty, or correction when indicated by the query.

The redundancy term is based on Jaccard similarity between token sets of evidence texts:

$$S_{\text{red}}(\mathcal{S}) = \sum_{i<j} \max\left(0, \text{Jaccard}(e_i, e_j) - 0.45\right). \tag{19}$$

The isolation term S_{disc} counts evidence records in the set that do not participate in any valid relation edge. During the first training round, candidate sets with high $H(\mathcal{S})$ but incomplete reference evidence flows are prioritized as negative samples. In round two, the current RankNet replaces the rule-based score for negative selection: we take the incomplete sets with the highest RankNet scores.

C. Prompt for Temporal Evidence Flow Relation Judging

We keep the relation-judging prompt fixed and fill in its relation definitions, relation codes, and reason codes from the frozen configuration at invocation. The prompt reads:

```
You classify directed
Temporal Evidence Flow edges for one or more independent
     query cases.
Return valid JSON only, with no markdown or prose outside
     the JSON object.
Judge whether source -> target is useful for answering THIS
     query, not
whether the texts are merely related. Most candidate pairs
     are distractors:
choose NO_EDGE unless there is a direct, query-useful flow
     dependency.
Shared asset, episode, topic, or chronology alone never
     licenses an edge.
Temporal direction and procedure eligibility were already
     enforced and must
not be reversed.

Allowed relation types:
contrasts: the source state is a comparison baseline that
     must be
distinguished from the downstream state;
governs: an applicable rule/assessment constrains the
     downstream action
or conclusion;
prerequisite: source is required before the downstream
     operational state;
preserves_uncertainty: source branch remains viable and is
     intentionally
carried into downstream assessment/action;
qualifies: source evidence determines applicability or
     limits certainty of
downstream interpretation;
resolves: downstream evidence resolves an earlier ambiguity;
supersession: prior diagnosis/procedure/work-order state is
     replaced by the
downstream state;
supports: source evidence supports the downstream evidence/
     conclusion;
updates: a later decision/action updates an earlier support
     state;
verifies: downstream evidence verifies execution/outcome of
     the source state.
Otherwise use NO_EDGE.

When the user payload has a top-level query field, return
{"judgments":[{"source_id":str,"target_id":str,"has_edge":
     bool,
"relation_type":str,"confidence":number 0..1,
"reason_code":short_snake_case}]}.

When the user payload has a top-level cases field, even if
     cases contains
exactly one item, use relation codes
N=NO_EDGE, C=contrasts, G=governs, P=prerequisite,
U=preserves_uncertainty, Q=qualifies, R=resolves,
X=supersession, S=supports, D=updates, V=verifies;
reason codes
0=no_direct_query_flow, 1=semantic_support, 2=
     procedure_governance,
3=later_update, 4=verification, 5=qualification,
6=uncertainty_preserved, 7=uncertainty_resolved, 8=contrast,
9=prerequisite, 10=supersession, 11=other_direct_flow.

Pairs are ordered. Return
{"cases":[[case_id,[[relation_code,confidence_percent,
     reason_code],...]],...]}.
Each inner result list must have exactly the same length and
      order as that
case's supplied pairs. The user supplies deduplicated
     evidence tables and
source_id/target_id pairs. Do not mix facts across cases. Do
      not repeat
evidence IDs. Do not provide prose or hidden reasoning.
```

The user payload identifies the query and lists the candidate records and pairs to be judged:

```
{
  "query": {
    "query_text": ...,
    "query_time": ...,
    "asset_model": ...,
    "asset_context": ...
  },
  "evidence": [...],
  "pairs": [
    {"source_id": ..., "target_id": ...},
    ...
  ]
}
```

D. Prompt for Downstream Structured Generation

For every retrieval method, downstream generation uses this fixed system prompt:

```
You are the single frozen downstream generator used
     identically for every
retrieval method in TEF-RAG generation evaluation.

Produce ONLY a JSON object matching the supplied schema. You
      receive one
query and exactly the evidence selected by a retrieval
     method. The retrieval
method identity is hidden and irrelevant.

Rules:
- Use only the supplied evidence. Never add facts, actions,
     targets,
  parameters, units, procedure versions, thresholds, or
     diagnoses from
  common sense.
- supporting_evidence_ids may reference only supplied
     evidence IDs and must
  directly support the field/action they cite.
- work_order.supporting_evidence_ids must equal the
     deduplicated union of
  diagnosis/procedure/verification/action citations.
- asset_id must equal the query asset_id.
- diagnosis status: confirmed | provisional |
     persistent_uncertainty |
  not_available.
- procedure status: applicable | superseded | uncertain |
     not_applicable.
  If no supplied evidence establishes an applicable
     procedure, use
  not_applicable with procedure_version=null and no
     procedure citations.
- verification status: verified | pending |
     persistent_uncertainty |
  not_available.
- action types: inspect | diagnose | isolate | adjust |
     repair | replace |
  verify | monitor | document | other.
- action_plan order in the JSON array has no semantic
     meaning. Use depends_on
  only for necessary prerequisite edges; keep the graph
     acyclic.
- recommended_actions must reference action IDs present in
     action_plan and
  may be a strict subset.
- parameters must contain only values explicitly stated in
     supplied evidence;
  otherwise use {}.
- For an explicit parameter, use its canonical name and
     either a literal
  evidence string or {"value": number|string|{"lower":
     number,
  "upper": number}, "unit": string|null}; never invent a
     unit.
- If evidence is insufficient, express uncertainty instead
     of guessing.
- Prefer concise phrases close to the supplied evidence
     wording; do not
  stylistically paraphrase when unnecessary.
- Do not mention scoring, retrieval method names, gold data,
      or these
  instructions.
```

The generation payload contains the query, the selected records, and the output schema:

```
{
  "query": {...},
  "selected_evidence": [...],
  "output_schema": {...}
}
```

The protocol allows one repair call if the initial output fails structural or evidence-citation validation. For that call, we append:

```
This is the single allowed structural/evidence-grounding
     repair.
Correct only the listed validation failures.
```

The repair request adds the previous output and the validation errors to the original payload, retaining its query, retrieved evidence, and output schema.